The CoMeT User Manual

June 29, 2023

This guide builds on top of the *HotSniper* user manual available here.

1 Write a Custom DVFS Policy

This guide provides step-by-step instructions to implement a new policy. We take the Linux ondemand governor as an example of a DVFS policy. To showcase the use of integrated thermal feedback, we implement a reactive Dynamic Thermal Management (DTM) that responds to peak core temperatures.

1. Create the files `dvfsOndemand.cc` and `dvfsOndemand.h` in the folder `common/scheduler/policies`.

Our policy requires access to the performance counters, as well as several configurable parameters. We provide them to the policy in the constructor. We implement the `DVFSPolicy` interface that requires a function `getFrequencies` as indicated below. Finally, we define some class member variables to hold some internal state.

2. Change the code of `dvfsOndemand.h` as listed in Appendix A.1.

Next, we provide the actual implementation of our policy. The policy is queried periodically (method `getFrequencies`). This function returns a new vector of per-core frequencies to be used.

3. Change the code of `dvfsOndemand.cc` as listed in Appendix A.2.

We defined several parameters for our policy that we would like to expose in the configuration files. In the next step, we add them there and provide values. Using namespaces (e.g.: `[scheduler/open/dvfs/ondemand]`) is optional but helps structure the code better.

4. Add the following code to `config/base.cfg`:

```
[scheduler/open/dvfs/ondemand]
up_threshold = 0.7
down_threshold = 0.3
dtm_cricital_temperature = 80
dtm_recovered_temperature = 78
```

We are now ready to instantiate our policy.

5. Include our new header file in `common/scheduler/scheduler_open.cc`:

```
#include "policies/dvfsOndemand.h"
```

6. Extend the method `initDVFSPolicy` in `common/scheduler/scheduler_open.cc`:

```
} else if (policyName == "ondemand") {
        float upThreshold = Sim()->getCfg()->getFloat(
            "scheduler/open/dvfs/ondemand/up_threshold");
        float downThreshold = Sim()->getCfg()->getFloat(
            "scheduler/open/dvfs/ondemand/down_threshold");
        float dtmCriticalTemperature = Sim()->getCfg()->getFloat(
            "scheduler/open/dvfs/ondemand/dtm_cricital_temperature");
        float dtmRecoveredTemperature = Sim()->getCfg()->getFloat(
            "scheduler/open/dvfs/ondemand/dtm_recovered_temperature");
        dvfsPolicy = new DVFSOndemand(
                performanceCounters,
                coreRows,
                coreColumns,
                minFrequency,
                maxFrequency,
                frequencyStepSize,
                upThreshold,
                downThreshold,
                dtmCriticalTemperature,
                dtmRecoveredTemperature
        );
} //else if (policyName ="XYZ") {... } //Place to instantiate a new DVFS logic.
    Implementation is put in "policies" package.
```

This code checks whether the policy with the name `ondemand` shall be used. If this is the case, it first reads the four configuration options and then instantiates our policy.

Finally, we need to select our new policy in the config.

7. Perform the following change in `config/base.cfg`:

```
[scheduler/open/dvfs]
logic = ondemand
```

Finally, rebuild CoMeT.

8. Execute `make` in the base folder.

2 Write a Custom Task Mapping and Migration Policy

This section describes how to create a thermal-aware task mapping and migration policy. When a new application arrives, we want to map it to the coldest cores. When cores get too hot, we want to migrate the threads away from them to the coldest cores. This policy also serves as an example of how to build a policy that uses several knobs (mapping and migration) simultaneously.

1. Create the files `coldestCore.cc` and `coldestCore.h` in the folder `common/scheduler/policies`.

As this policy should perform both task mapping and migration, we implement both interfaces `MappingPolicy` and `MigrationPolicy`. If a policy should also do DVFS, additionally implement the `DVFSPolicy` as in Section 1. The open scheduler calls the `map` function whenever a new application needs to be mapped, and periodically calls the `migrate` function.

2. Change the code of `coldestCore.h` as listed in Appendix A.3.

We now implement the logic of these functions. `map` needs to return the list of cores for this application (in the one-thread-per-core model). `migration` can return an arbitrarily long list of

thread migrations to be performed. Each migration needs to specify the source and target core of the migration. The `swap` attribute is used to indicate that both source and target core are used, and the corresponding threads should be swapped. While this information would also be present in the scheduler, the requirement to explicitly mark swap operations serves to detect bugs in the migration policy early.

3. Change the code of `coldestCore.cc` as listed in Appendix A.4.

Our policy migrates threads away from hot cores when the temperature has exceeded a certain threshold. This threshold is passed to the constructor. We want to have this value easily configurable. First, we add the value to a new config section for our policy. Notice that adding a new section is optional, but maintains the readability of the config.

4. Add the following code to `config/base.cfg`:

```
[scheduler/open/migration/coldestCore]
criticalTemperature = 80
```

We are now ready to instantiate our policy.

5. Include our new header file in `common/scheduler/scheduler_open.cc`:

```
#include "policies/coldestCore.h"
```

As we perform both mapping and migration, we need to add our policy to two places.

6. Extend the method `initMappingPolicy` in `common/scheduler/scheduler_open.cc`:

```
} else if (policyName == "coldestCore") {
        float criticalTemperature = Sim()->getCfg()->getFloat(
                "scheduler/open/migration/coldestCore/criticalTemperature");
        mappingPolicy = new ColdestCore(performanceCounters, coreRows,
            coreColumns, criticalTemperature);
} //else if (policyName ="XYZ") {... } //Place to instantiate a new mapping
    logic. Implementation is put in "policies" package.
```

7. Extend the method `initMigrationPolicy` in `common/scheduler/scheduler_open.cc`:

```
} else if (policyName == "coldestCore") {
        float criticalTemperature = Sim()->getCfg()->getFloat(
                "scheduler/open/migration/coldestCore/criticalTemperature");
        migrationPolicy = new ColdestCore(performanceCounters, coreRows,
            coreColumns, criticalTemperature);
} //else if (policyName ="XYZ") {... } //Place to instantiate a new migration
    logic. Implementation is put in "policies" package.
```

Now, we are ready to select our new policy for both mapping and migration.

8. Perform the following changes in `config/base.cfg`:

```
[scheduler/open]
logic = coldestCore

[scheduler/open/migration]
logic = coldestCore
```

Finally, rebuild CoMeT.

9. Execute `make` in the base folder.

3 Write a Custom Memory Bank DTM Policy

The basic functionality of the memory DTM policy can be seen in `drampolicy.h` in `common/scheduler/policies/`. The open scheduler will call `std::map<int,int> getNewBankModes()`, which will return a map where the key is the bank number, and the value is either 0 or 1, where 0 means low power mode and 1 means normal power mode. A memory DTM policy will be an extension of this file.

1. Create the files `dramLowpower.cc` and `dramLowpower.h` in `common/scheduler/policies`.

   We will give our memory DTM access to the performance counters to retrieve the bank temperatures, as well as some parameters that are required for the policy, such as the number of banks and the temperature thresholds. These are defined in the header file.

2. Copy the code from appendix A.5 to `dramLowpower.h`.

   The next step will be to implement the policy. After every epoch, the policy's function `getNewBankModes` is called, which will return a map with the power mode per memory bank.

3. Copy the code from appendix A.6 to `dramLowpower.cc`.

   The parameters for the policy will have to be provided through the configuration files

4. In `config/base.cfg`, add the following two lines:

   ```
   [scheduler/open/dram/dtm]
   dtm_critical_temperature = 68
   dtm_recovered_temperature = 65
   ```

   There should already be two lines under `[scheduler/open/dram/dtm]`, namely

   ```
   bank_mode_trace_file = bank_mode.trace
   full_bank_mode_trace_file = full_bank_mode.trace
   ```

   If they are not there, you must include them now.

   Let us now add our policy.

5. Include the header file in `common/scheduler/scheduler_open.cc`:

   ```
   #include "policies/dramLowpower.h"
   ```

6. Extend the method `initDramPolicy` in `common/scheduler/scheduler_open.cc`:

   ```
   } else if (policyName == "lowpower") {

           float dtmCriticalTemperature = Sim()->getCfg()->getFloat("scheduler/
               open/dram/dtm/dtm_critical_temperature");
           float dtmRecoveredTemperature = Sim()->getCfg()->getFloat("scheduler
               /open/dram/dtm/dtm_recovered_temperature");
           dramPolicy = new DramLowpower(
                   performanceCounters,
                   numberOfBanks,
                   dtmCriticalTemperature,
                   dtmRecoveredTemperature
           );
   } // else if (policyName ="XYZ") {... } // Instantiate a new memory DTM
       policy here. Implementation is put in "policies" package.
   ```

This code checks if the name of the memory DTM policy is equal to `lowpower`, and if so, reads the values from `config/base.cfg` which we provided in an earlier step. We will now add this name to the same configuration file.

7. Add the following to `config/base.cfg`:

```
[scheduler/open/dram]
dtm = lowpower
```

Finally, rebuild CoMeT:

8. You can now run `make` in the base folder.

Keep the following things in mind when experimenting with memory DTM. Make sure that the dram performance model is set to `normal`:

```
[perf_model/dram]
type = normal
```

Also make sure to use the open scheduler:

```
[scheduler]
type = open
```

The epoch of the memory DTM is set with:

```
[scheduler/open/dram]
dram_epoch = 1000000
```

The latency of memory banks in low power mode is set with:

```
[perf_model/dram]
latency_lowpower = 6000
```

The power fractions of banks in low power mode are set with:

```
[perf_model/dram/lowpower]
lpm_dynamic_power = 0.01
lpm_leakage_power = 0.1
```

A   Code

A.1   `dvfsOndemand.h`

```
/**
 * This header implements the ondemand governor with DTM.
 * The ondemand governor implementation is based on
 * Pallipadi, Venkatesh, and Alexey Starikovskiy.
 * "The ondemand governor."
 * Proceedings of the Linux Symposium. Vol. 2. No. 00216. 2006.
 */

#ifndef __DVFS_ONDEMAND_H
#define __DVFS_ONDEMAND_H

#include <vector>
#include "dvfspolicy.h"
#include "performance_counters.h"
```

```cpp
class DVFSOndemand : public DVFSPolicy {
public:
    DVFSOndemand(
        const PerformanceCounters *performanceCounters,
        int coreRows,
        int coreColumns,
        int minFrequency,
        int maxFrequency,
        int frequencyStepSize,
        float upThreshold,
        float downThreshold,
        float dtmCriticalTemperature,
        float dtmRecoveredTemperature);
    virtual std::vector<int> getFrequencies(
        const std::vector<int> &oldFrequencies,
        const std::vector<bool> &activeCores);

private:
    const PerformanceCounters *performanceCounters;

    unsigned int coreRows;
    unsigned int coreColumns;
    int minFrequency;
    int maxFrequency;
    int frequencyStepSize;
    float upThreshold;
    float downThreshold;
    float dtmCriticalTemperature;
    float dtmRecoveredTemperature;

    bool in_throttle_mode = false;
    bool throttle();
};

#endif
```

A.2 dvfsOndemand.cc

```cpp
#include "dvfsOndemand.h"
#include <iomanip>
#include <iostream>

using namespace std;

DVFSOndemand::DVFSOndemand(
        const PerformanceCounters *performanceCounters,
        int coreRows,
        int coreColumns,
        int minFrequency,
        int maxFrequency,
        int frequencyStepSize,
        float upThreshold,
        float downThreshold,
```

```cpp
        float dtmCriticalTemperature,
        float dtmRecoveredTemperature)
    : performanceCounters(performanceCounters),
      coreRows(coreRows),
      coreColumns(coreColumns),
      minFrequency(minFrequency),
      maxFrequency(maxFrequency),
      frequencyStepSize(frequencyStepSize),
      upThreshold(upThreshold),
      downThreshold(downThreshold),
      dtmCriticalTemperature(dtmCriticalTemperature),
      dtmRecoveredTemperature(dtmRecoveredTemperature) {

}

std::vector<int> DVFSOndemand::getFrequencies(
        const std::vector<int> &oldFrequencies,
        const std::vector<bool> &activeCores) {
    if (throttle()) {
        std::vector<int> minFrequencies(coreRows * coreColumns, minFrequency);
        cout << "[Scheduler][ondemand-DTM]: in throttle mode -> return min.
            frequencies" << endl;
        return minFrequencies;
    } else {
        std::vector<int> frequencies(coreRows * coreColumns);

        for (unsigned int coreCounter = 0; coreCounter < coreRows * coreColumns;
            coreCounter++) {
            if (activeCores.at(coreCounter)) {
                float power = performanceCounters->getPowerOfCore(coreCounter);
                float temperature = performanceCounters->getTemperatureOfCore(
                    coreCounter);
                int frequency = oldFrequencies.at(coreCounter);
                float utilization = performanceCounters->getUtilizationOfCore(
                    coreCounter);

                cout << "[Scheduler][ondemand]: Core " << setw(2) << coreCounter
                    << ":";
                cout << " P=" << fixed << setprecision(3) << power << " W";
                cout << " f=" << frequency << " MHz";
                cout << " T=" << fixed << setprecision(1) << temperature << " C";
                    // avoid the little circle symbol, it is not ASCII
                cout << " utilization=" << fixed << setprecision(3) << utilization
                     << endl;

                // use same period for upscaling and downscaling as described
                // in "The ondemand governor."
                if (utilization > upThreshold) {
                    cout << "[Scheduler][ondemand]: utilization > upThreshold";
                    if (frequency == maxFrequency) {
                        cout << " but already at max frequency" << endl;
                    } else {
                        cout << " -> go to max frequency" << endl;
                        frequency = maxFrequency;
```

```cpp
                }
            } else if (utilization < downThreshold) {
                cout << "[Scheduler][ondemand]: utilization < downThreshold";
                if (frequency == minFrequency) {
                    cout << " but already at min frequency" << endl;
                } else {
                    cout << " -> lower frequency" << endl;
                    frequency = frequency * 80 / 100;
                    frequency = (frequency / frequencyStepSize) *
                        frequencyStepSize; // round
                    if (frequency < minFrequency) {
                        frequency = minFrequency;
                    }
                }
            }

            frequencies.at(coreCounter) = frequency;
        } else {
            frequencies.at(coreCounter) = minFrequency;
        }
    }
    return frequencies;
    }
}

bool DVFSOndemand::throttle() {
    if (performanceCounters->getPeakTemperature() > dtmCriticalTemperature) {
        if (!in_throttle_mode) {
            cout << "[Scheduler][ondemand-DTM]: detected thermal violation" <<
                endl;
        }
        in_throttle_mode = true;
    } else if (performanceCounters->getPeakTemperature() <
        dtmRecoveredTemperature) {
        if (in_throttle_mode) {
            cout << "[Scheduler][ondemand-DTM]: thermal violation ended" << endl;
        }
        in_throttle_mode = false;
    }
    return in_throttle_mode;
}
```

A.3   coldestCore.h

```cpp
/**
 * This header implements a policy that maps new applications to the coldest
    core
 * and migrates threads from hot cores to the coldest cores.
 */

#ifndef __COLDESTCORE_H
#define __COLDESTCORE_H
```

```cpp
#include <vector>
#include "mappingpolicy.h"
#include "migrationpolicy.h"
#include "performance_counters.h"

class ColdestCore : public MappingPolicy, public MigrationPolicy {
public:
    ColdestCore(
        const PerformanceCounters *performanceCounters,
        int coreRows,
        int coreColumns,
        float criticalTemperature);
    virtual std::vector<int> map(
        String taskName,
        int taskCoreRequirement,
        const std::vector<bool> &availableCores,
        const std::vector<bool> &activeCores);
    virtual std::vector<migration> migrate(
        SubsecondTime time,
        const std::vector<int> &taskIds,
        const std::vector<bool> &activeCores);

private:
    const PerformanceCounters *performanceCounters;

    unsigned int coreRows;
    unsigned int coreColumns;
    float criticalTemperature;

    int getColdestCore(const std::vector<bool> &availableCores);
    void logTemperatures(const std::vector<bool> &availableCores);
};

#endif
```

A.4 `coldestCore.cc`

```cpp
#include "coldestCore.h"

#include <iomanip>

using namespace std;

ColdestCore::ColdestCore(
        const PerformanceCounters *performanceCounters,
        int coreRows,
        int coreColumns,
        float criticalTemperature)
    : performanceCounters(performanceCounters),
      coreRows(coreRows),
      coreColumns(coreColumns),
      criticalTemperature(criticalTemperature) {

}

std::vector<int> ColdestCore::map(
        String taskName,
        int taskCoreRequirement,
        const std::vector<bool> &availableCoresRO,
        const std::vector<bool> &activeCores) {

    std::vector<bool> availableCores(availableCoresRO);

    std::vector<int> cores;

    logTemperatures(availableCores);

    for (; taskCoreRequirement > 0; taskCoreRequirement--) {
        int coldestCore = getColdestCore(availableCores);

        if (coldestCore == -1) {
            // not enough free cores
            std::vector<int> empty;
            return empty;
        } else {
            cores.push_back(coldestCore);
            availableCores.at(coldestCore) = false;
        }
    }

    return cores;
}

std::vector<migration> ColdestCore::migrate(
        SubsecondTime time,
        const std::vector<int> &taskIds,
        const std::vector<bool> &activeCores) {

    std::vector<migration> migrations;
```

```cpp
        std::vector<bool> availableCores(coreRows * coreColumns);
        for (int c = 0; c < coreRows * coreColumns; c++) {
            availableCores.at(c) = taskIds.at(c) == -1;
        }

        for (int c = 0; c < coreRows * coreColumns; c++) {
            if (activeCores.at(c)) {
                float temperature = performanceCounters->getTemperatureOfCore(c);
                if (temperature > criticalTemperature) {
                    cout << "[Scheduler][coldestCore-migrate]: core" << c << " too hot
                        (";
                    cout << fixed << setprecision(1) << temperature << ") -> migrate";
                    logTemperatures(availableCores);

                    int targetCore = getColdestCore(availableCores);

                    if (targetCore == -1) {
                        cout << "[Scheduler][coldestCore-migrate]: no target core
                            found, cannot migrate" << endl;
                    } else {
                        migration m;
                        m.fromCore = c;
                        m.toCore = targetCore;
                        m.swap = false;
                        migrations.push_back(m);
                        availableCores.at(targetCore) = false;
                    }
                }
            }
        }

        return migrations;
}

int ColdestCore::getColdestCore(const std::vector<bool> &availableCores) {
        int coldestCore = -1;
        float coldestTemperature = 0;
        // iterate all cores to find coldest
        for (int c = 0; c < coreRows * coreColumns; c++) {
            if (availableCores.at(c)) {
                float temperature = performanceCounters->getTemperatureOfCore(c);
                if ((coldestCore == -1) || (temperature < coldestTemperature)) {
                    coldestCore = c;
                    coldestTemperature = temperature;
                }
            }
        }
        return coldestCore;
}

void ColdestCore::logTemperatures(const std::vector<bool> &availableCores) {
    cout << "[Scheduler][coldestCore-map]: temperatures of available cores:" <<
        endl;
```

```cpp
    for (int y = 0; y < coreRows; y++) {
        for (int x = 0; x < coreColumns; x++) {
            if (x > 0) {
                cout << " ";
            }
            int coreId = y * coreColumns + x;

            if (!availableCores.at(coreId)) {
                cout << " - ";
            } else {
                float temperature = performanceCounters->getTemperatureOfCore(
                    coreId);
                cout << fixed << setprecision(1) << temperature;
            }
        }
        cout << endl;
    }
}
```

A.5 `dramLowpower.h`

```cpp
/**
 * This header implements memory DTM using a low power mode.
 */

#ifndef __DRAM_LOWPOWER_H
#define __DRAM_LOWPOWER_H

#include <map>

#include "drampolicy.h"
#include "performance_counters.h"

class DramLowpower : public DramPolicy {
public:
    DramLowpower(
        const PerformanceCounters *performanceCounters,
        int numberOfBanks,
        float dtmCriticalTemperature,
        float dtmRecoveredTemperature);
    virtual std::map<int,int> getNewBankModes(std::map<int,int> old_bank_modes);

private:
    const PerformanceCounters *performanceCounters;

    unsigned int numberOfBanks;
    float dtmCriticalTemperature;
    float dtmRecoveredTemperature;
};

#endif
```

A.6 `dramLowpower.cc`

```cpp
#include "dramLowpower.h"
#include <iomanip>
#include <iostream>
#include <map>

using namespace std;

DramLowpower::DramLowpower(
        const PerformanceCounters *performanceCounters,
        int numberOfBanks,
        float dtmCriticalTemperature,
        float dtmRecoveredTemperature)
    : performanceCounters(performanceCounters),
      numberOfBanks(numberOfBanks),
      dtmCriticalTemperature(dtmCriticalTemperature),
      dtmRecoveredTemperature(dtmRecoveredTemperature) {

}

/*
Return the new memory modes, based on current temperatures.
*/
std::map<int,int> DramLowpower::getNewBankModes(std::map<int, int>
    old_bank_modes) {

    cout << "in DramLowpower::getNewBankModes\n";
    std::map<int,int> new_bank_mode_map;
    for (int i = 0; i < numberOfBanks; i++)
    {
        if (old_bank_modes[i] == 0) // if the memory was already in low power
            mode
        {
            if (performanceCounters->getTemperatureOfBank(i) <
                dtmRecoveredTemperature) // temp dropped below recovery
                temperature
            {
                cout << "[Scheduler][dram-DTM]: thermal violation ended for bank "
                    << i << endl;
                new_bank_mode_map[i] = 1;
            }
            else
            {
                new_bank_mode_map[i] = 0;
            }
        }
        else // if the memory was not in low power mode
        {
            if (performanceCounters->getTemperatureOfBank(i) >
                dtmCriticalTemperature) // temp is above critical temperature
            {
                cout << "[Scheduler][dram-DTM]: thermal violation detected for
                    bank " << i << endl;
```

```
                new_bank_mode_map[i] = 0;
            }
            else
            {
                new_bank_mode_map[i] = 1;
            }
        }
    }
    return new_bank_mode_map;
}
```